6.824 2021 Midterm Exam

This is a 90-minute exam. Please write down any assumptions you make. You are allowed to consult 6.824 papers, notes, and lab code. If you have questions, please ask them on Piazza in a private post to the staff, or ask in the Zoom chat if you are taking the exam during lecture time. Please don't collaborate or discuss the exam with anyone.

In particular, please do not discuss the contents of the exam with anyone at all besides the course staff until at least 3:00 PM EDT on Friday, April 2nd, to allow students in other timezones to take the test.

MapReduce

In lab 1, a worker in the MapReduce library doesn't hand out any reduce tasks until all map tasks have completed and have produced all their intermediate files. The implementation described in the "MapReduce" paper by Dean and Ghemawat allows reduce workers to start reading intermediate files *before* all map tasks have completed.

A. Briefly describe the advantage of allowing reduce workers to start reading intermediate files early.

B. Describe how you would modify your lab implementation to allow reading of intermediate files early, as Google's MapReduce library does, while maintaining correctness. (You don't have to write code, but sketch out a design in words. Keep your answer brief.)

MapReduce (Lab)

Zara Zoom has implemented MapReduce for 6.824 Lab 1. Their worker code loops, and asks the coordinator for tasks using a `GetTask` RPC call. Their coordinator code to handle `GetTask` RPCs and assign tasks to workers is as follows:

```
func (c *Coordinator) HandleGetTask(args *GetTaskArgs, reply *GetTaskReply) error {
    c.mu.Lock()
    defer c.mu.Unlock()

    for {
        // iterate through tasks that aren't done
        for i := c.numCompletedMapTasks; i < c.numTotalMapTasks; i++ {
            if c.mapStarted[i] == false || time.Since(c.mapWhen[i]).Seconds() > 5) {
                c.mapStarted[i] = true
                c.mapWhen[i] = time.Now()

                reply.TaskType = MapTask
                reply.TaskNum = i
                // fill in rest of GetTaskReply fields

                // remember that we assigned a task
                return nil
            }
        }

        // if we've completed all map tasks, break out of the loop that
        // assigns map tasks
        if c.numCompletedMapTasks >= c.numTotalMapTasks {
            break
        }

        // if we got here, all map tasks have been started,
        // but some have not finished.
        // wait for either a map task to complete,
        // or for some timeout to occur to wake us up.
        c.cond.Wait()
    }

    // all maps have completed. move on to reduce phase.
    ...
    // when all reduces have completed, tell worker to exit with a
    // DoneTask.
    reply.TaskType = DoneTask
    return nil
}
```

Zara also has their workers send `FinishedTask` RPCs to tell the coordinator when a particular map task has finished.

```
func (c *Coordinator) HandleFinishedTask(args *FinishedTaskArgs, reply *FinishedTaskReply) error {
    c.mu.Lock()
    defer c.mu.Unlock()

    if args.Type == MapTask {
        if c.numCompletedMapTasks == c.numTotalMapTasks {
            fmt.Printf("\map task assigned twice, already done\n\n")
        }
        c.numCompletedMapTasks += 1
    } else if args.Type == ReduceTask {
        ...
    }

    // wake up waiting GetTask handlers
    c.cond.Broadcast()
    return nil
}
```

You can assume that all code not shown is correct.

When testing their implementation, Zara notices that a job with map tasks numbered 0, 1, and 2 often fails.

A. In some scenarios, some reduce workers cannot open intermediate file `mr-1-x` (where x is the Reduce task number). Explain a scenario that could cause this to occur.

B. Briefly describe what Zara should do to fix this problem.

GFS

Consider the paper "The Google File System" by Ghemawat et al.

A. The GFS file system doesn't guarantee linearizability. Give an example of how the lack of linearizability complicates writing programs with GFS. That is, describe a feature that one must implement when using GFS that one wouldn't have to implement if GFS would have provided linearizability. Briefly explain your answers.

B. Briefly explain why append operations are not allowed to span chunk boundaries.

VMware FT

In VM-FT (as described by in the paper "Fault-Tolerant Virtual Machines" by Scales et al.), the backup lags behind the primary to ensure, for example, that network interrupts are delivered at exactly the same instruction at the backup as on the primary. Briefly explain what could go wrong if the backup delivered interrupts at a different instruction than the primary?

Raft

Consider figure 7 of the Raft paper "In search of Understandable Consensus Algorithms (Extended Version)". The last follower (at the bottom of the figure) has term 3 in its last index 11. Could there have been a scenario under which that follower's log had entries in indexes 12, 13, and 14 in figure 7? What terms could those entries have? (Briefly explain your answer.)

Raft (Lab)

Rob Remote managed to get a correct Raft implementation through Lab 2C (persistence). He is interested in analyzing his correct system. Specifically, Rob wants to know how many times a particular value passed into `Start()` is sent over any of the servers' `applyCh`. Suppose he has 5 servers.

After running for a while, the tester calls `Start()` with a command value 2 and returns `isLeader=true` and `index=100`.

Determine all possible values for the **number of times** command 2 is sent through an `applyCh` in the following cases (i.e., if a command is delivered once on each peer's `applyCh`, then the total number of times is 5). Briefly explain your answers.

A. No failures of any kind

B. Network partitions possible (but no crashes)

C. Crashes possible (but no network partitions)

Raft (Persistence)

While implementing Raft, Connie Consensus decided that instead of persisting the entire log, she would persist only committed entries. Connie claims that uncommitted entries could have been lost because of crashes or network issues so it's safe to ignore them until they've been committed. Does this change affect Raft's correctness? If so, provide a sequence of events that would lead to unsafe behavior. You can assume that the rest of Connie's implementation is correct.

Zookeeper

Consider Zookeeper, as described in the paper "ZooKeeper: Wait-free coordination for Internet-scale systems".

A. Zookeeper's read throughput scales with the number of servers (e.g., see the red line in Figure 5). The Raft paper describes an optimization for read-only operations. Would that optimization allow Raft to scale read throughput with the number of peers as in Zookeeper? (Briefly explain your answer.)

B. ZooKeeper performs all operations in FIFO client order. Consider the following sequence of operations starting in a state where `f` contains 0 (time goes down):

```
 Client 1:        Client 2:
 f contains 0
 delete("ready")
 write 1 to f
 create("ready")
                  exists("ready")
                  read f
```

For this example, briefly explain what could go wrong if ZooKeeper didn't guarantee FIFO client order. In particular, what values could the read of `f` by client 2 return?

6.824

A. Which papers/lectures should we omit in future 6.824 years, because they are not useful or are too hard to understand?

[] MapReduce
[] Q&A Lecture MapReduce
[] GFS
[] VMware FT
[] Raft
[] ZooKeeper
[] Q&A Lecture Raft

B. Which papers/lectures did you find most useful?

[] MapReduce
[] Q&A Lecture MapReduce
[] GFS
[] VMware FT
[] Raft
[] Q&A Lecture Raft
[] ZooKeeper

C. What should we change about the course to make it better?